
SC4064 Proposal

Aryan Jain* Fanyi Pu* Ze Hong Maxwell Au*

School of Computer Science and Engineering

Nanyang Technological University, Singapore

{ARYAN017, FPU001, MAU002}@e.ntu.edu.sg

Abstract

General Matrix–Matrix Multiplication (GEMM) is the computational backbone of modern deep learning. This project bridges *kernel-level optimization* and *system-level parallelism* by (i) implementing and progressively optimizing CUDA GEMM kernels on a single GPU, and (ii) extending to a multi-GPU Tensor Parallel linear layer with NCCL-based communication. We systematically evaluate the compute–communication trade-off, strong/weak scaling behavior, and the impact of kernel-level optimization on distributed scaling efficiency.

1 Introduction

General Matrix–Matrix Multiplication (GEMM) underpins virtually all compute-intensive operations in deep learning [Jia et al., 2018]. In transformer architectures [Vaswani et al., 2017], multi-head attention and feedforward layers reduce to large matrix multiplications, making GEMM performance a first-order concern. As model sizes grow beyond the memory capacity of a single accelerator, **Tensor Parallelism** [Shoeybi et al., 2019] has become standard for distributing parameters across GPUs, where inter-GPU communication becomes a critical throughput bottleneck.

Despite vendor-optimized libraries such as cuBLAS [NVIDIA, 2024], understanding the interplay between kernel-level compute efficiency and system-level communication overhead remains essential. This project aims to: (1) systematically optimize GEMM on a single GPU using CUDA; (2) implement a multi-GPU Tensor Parallel linear layer (forward and backward) using NCCL; and (3) quantitatively analyze how kernel optimization interacts with distributed scaling efficiency.

2 Proposed methodology

Single-GPU GEMM optimization. We implement GEMM from scratch in CUDA and progressively optimize: (1) **naive implementation** using global memory with coalesced access; (2) **tiled GEMM with shared memory** to exploit data locality; (3) **register blocking and loop unrolling** to increase arithmetic intensity; (4) **Tensor Core acceleration** (optional) via WMMA intrinsics for mixed-precision computation. At each stage, we profile using NVIDIA Nsight Compute to measure occupancy, bandwidth utilization, and throughput (GFLOP/s), benchmarking against cuBLAS. We employ *roofline model* analysis [Williams et al., 2009] to classify each variant as compute-bound or memory-bound.

Multi-GPU Tensor Parallel forward pass. We implement column-wise Tensor Parallelism for a linear layer. Given $Y = XW$, the weight matrix is partitioned column-wise across p GPUs: $W = [W_1, \dots, W_p]$, where GPU i computes $Y_i = XW_i$. Outputs are gathered via NCCL AllGather. We investigate communication–computation overlap using CUDA streams, the effect of partition granularity, and communication overhead as a function of GPU count.

*Equal contribution. Authors listed in alphabetical order.

Backward pass and gradient aggregation. Each GPU computes local gradients: $\partial\mathcal{L}/\partial W_i = X^\top(\partial\mathcal{L}/\partial Y_i)$ and $\partial\mathcal{L}/\partial X = \sum_{i=1}^p (\partial\mathcal{L}/\partial Y_i)W_i^\top$, where the latter requires an AllReduce across GPUs. We compare synchronous versus computation-overlapped reduction, revealing how communication cost dominates as local GEMM becomes highly optimized.

3 Evaluation plan

All experiments will be conducted on NVIDIA A100 GPUs.

Single-GPU performance. We measure throughput (GFLOP/s) across matrix sizes (512–8192) for each optimization stage, comparing against cuBLAS. Roofline analysis identifies bottleneck transitions between memory-bound and compute-bound regimes.

Strong and weak scaling. For strong scaling, we fix the problem size and increase GPUs (1, 2, 4, 8), measuring speedup and parallel efficiency. For weak scaling, we fix per-GPU workload and evaluate linear scalability.

Compute–communication trade-off. We measure the proportion of wall-clock time in NCCL communication versus GEMM computation, evaluating how kernel optimization shifts the bottleneck from compute to communication.

4 Expected contributions

This project delivers: (1) a profiling-driven study of CUDA GEMM optimization with roofline analysis; (2) a working multi-GPU Tensor Parallel linear layer (forward and backward) using NCCL; (3) quantitative analysis of compute–communication trade-offs and scaling behavior; and (4) insights into how kernel-level optimization interacts with distributed parallelism.

References

- Zhe Jia, Marco Maggioni, Benjamin Staiger, and Daniele P Scarpazza. Dissecting the NVIDIA Volta GPU architecture via microbenchmarking. *arXiv preprint arXiv:1804.06826*, 2018.
- NVIDIA. cuBLAS library. <https://developer.nvidia.com/cUBLAS>, 2024.
- Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009.