
High-Performance GEMM: From Kernel Optimization to Tensor Parallelism

Aryan Jain* Fanyi Pu* Ze Hong Maxwell Au*

School of Computer Science and Engineering

Nanyang Technological University, Singapore

{ARYAN017, FPU001, MAU002}@e.ntu.edu.sg

Abstract

General Matrix–Matrix Multiplication (GEMM) serves as the computational cornerstone of modern deep learning. This project bridges the gap between *kernel-level optimization* and *system-level parallelism* by (i) implementing and progressively optimizing CUDA GEMM kernels on a single GPU, and (ii) extending these primitives to a multi-GPU Tensor Parallel linear layer utilizing NCCL-based communication. We systematically evaluate the compute–communication trade-offs, analyze strong/weak scaling behaviors, and quantify the impact of kernel efficiency on distributed training performance.

1 Introduction

General Matrix Multiplications (GEMM) underpins virtually all compute-intensive operations in deep learning [Jia et al., 2018]. In transformer architectures [Vaswani et al., 2017], multi-head attention mechanisms and feed-forward networks rely heavily on large-scale matrix multiplications, establishing GEMM performance as a primary determinant of system throughput. As model parameters scale beyond the memory capacity of a single accelerator, **Tensor Parallelism** [Shoeybi et al., 2019] has emerged as a standard technique for distributing weights across GPUs, introducing inter-GPU communication as a critical bottleneck.

While vendor-optimized libraries such as cuBLAS [NVIDIA, 2024] provide highly tuned primitives, a deep understanding of the interplay between kernel-level compute efficiency and system-level communication overhead remains essential for system designers. This project aims to: (1) systematically optimize GEMM on a single GPU using CUDA; (2) implement a distributed Tensor Parallel linear layer (forward and backward) using NCCL; and (3) quantitatively analyze how local kernel optimization interacts with distributed scaling efficiency.

2 Proposed methodology

Single-GPU GEMM optimization. We develop a high-performance GEMM kernel from scratch in CUDA, following a progressive optimization roadmap: (1) **naive implementation** utilizing global memory with coalesced access; (2) **tiled GEMM with shared memory** to maximize data reuse and locality; (3) **register blocking and loop unrolling** to enhance arithmetic intensity; and (4) **Tensor Core acceleration** (optional) via WMMA intrinsics for mixed-precision performance. At each stage, we utilize NVIDIA Nsight Compute to profile occupancy, memory bandwidth, and compute throughput (GFLOP/s), benchmarking against cuBLAS. We employ *roofline model* analysis [Williams et al., 2009] to characterize performance bottlenecks and the transition from memory-bound to compute-bound regimes.

*Equal contribution. Authors listed in alphabetical order.

Multi-GPU Tensor Parallel forward pass. We implement column-wise Tensor Parallelism for a linear layer. Given $Y = XW$, the weight matrix is partitioned column-wise across p GPUs: $W = [W_1, \dots, W_p]$, such that GPU i computes $Y_i = XW_i$. Partial results are gathered via an NCCL AllGather operation. We investigate techniques for overlapping communication with computation using CUDA streams, analyze the impact of partition granularity, and quantify communication overhead as a function of GPU count.

Backward pass and gradient aggregation. During the backward pass, each GPU computes local gradients: $\partial\mathcal{L}/\partial W_i = X^\top(\partial\mathcal{L}/\partial Y_i)$ and $\partial\mathcal{L}/\partial X = \sum_{i=1}^p (\partial\mathcal{L}/\partial Y_i)W_i^\top$. The computation of $\partial\mathcal{L}/\partial X$ requires an AllReduce operation to aggregate gradients across GPUs. We compare synchronous execution against computation-overlapped reduction to demonstrate how communication latency becomes the dominant factor as local GEMM kernels are optimized.

3 Evaluation plan

We conduct all experiments on NVIDIA A100 GPUs.

Single-GPU performance. We measure throughput (GFLOP/s) across a range of matrix dimensions (512–8192) for each optimization stage, benchmarking against cuBLAS. Roofline analysis will be used to visualize the optimization trajectory.

Strong and weak scaling. For **strong scaling**, we fix the total problem size and increase the number of GPUs (1, 2, 4, 8) to measure speedup and parallel efficiency. For **weak scaling**, we fix the workload per GPU to evaluate linear scalability.

Compute–communication trade-off. We quantify the ratio of wall-clock time spent in NCCL communication versus GEMM computation. This metric allows us to evaluate how aggressive kernel optimization shifts the system bottleneck from computation to communication.

4 Expected contributions

Key contributions of this project include: (1) a comprehensive study of CUDA GEMM optimization driven by profiling and roofline analysis; (2) a functional implementation of a multi-GPU Tensor Parallel linear layer (forward and backward) using NCCL; (3) a quantitative analysis of compute–communication trade-offs and scaling characteristics; and (4) insights into the dependency of distributed scaling efficiency on underlying kernel performance.

References

- Zhe Jia, Marco Maggioni, Benjamin Staiger, and Daniele P Scarpazza. Dissecting the NVIDIA Volta GPU architecture via microbenchmarking. *arXiv preprint arXiv:1804.06826*, 2018.
- NVIDIA. cuBLAS library. <https://developer.nvidia.com/cUBLAS>, 2024.
- Mohammad Shoeybi, Mostafa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009.